

DHANANJAY SITAP

Data Engineer | AWS | PySpark | Databricks | SQL | GenAI

dhananjay.sitap.data@gmail.com | Maharashtra, India | +91 9325538419 | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Data Engineer with 1.6+ years of experience building production ETL/ELT pipelines on Databricks for a leading U.S. retail client (Target Corporation) using PySpark, SQL, and Delta Lake across a Medallion Architecture (Bronze–Silver–Gold). Enhanced the platform with a GenAI layer using Retrieval-Augmented Generation (RAG), ChromaDB, and LLM Tool Calling, enabling business users to query live retail data using natural language. Experienced in AWS (Amazon S3, Amazon Redshift, Amazon Athena), data quality engineering, production monitoring, Databricks Workflows, and CI/CD deployment using Azure DevOps. Seeking to grow into GenAI-focused Data Engineering roles.

PROFESSIONAL EXPERIENCE

Data Engineer

Jan 2025 - Present

2GBR Software Pvt. Ltd. | Client: Target Corporation

Technologies: Databricks, PySpark, Python, SQL, AWS S3, Redshift, Athena, Delta Lake, Azure DevOps, MySQL, ChromaDB, RAG, LLM Tool Calling, Power BI

Project Overview:

Target Corporation is one of the largest retail companies in the U.S. with over 1,800 stores and a strong omnichannel presence. The project focused on building a scalable data platform using Databricks, PySpark, and the Medallion Architecture to process 1.2–3 GB of daily retail data. The platform was further enhanced with a GenAI layer using RAG, ChromaDB, and LLM Tool Calling for natural language business analytics

- Engineered a retail data platform ingesting 1.2-3 GB of daily data from MySQL (Customer), REST APIs (Order, Payment), and JSON feeds (Product, Shipping) to replace fragmented, manually-generated business reporting.
- Built the Bronze layer on Databricks with PySpark, landing raw partitioned data in Amazon S3.
- Developed the Silver layer with null handling, duplicate removal, schema validation, data type conversion, and column standardization; routed invalid records to a rejected-records table.
- Produced business-ready Gold-layer Delta tables (Sales, Revenue, Customer, and Product Summaries) using ACID transactions and schema evolution.
- Loaded Gold-layer data into Amazon Redshift and authored analytical SQL reports using JOINS, GROUP BY, CTEs, and Window Functions; validated datasets in S3 with Amazon Athena.
- Enhanced the platform with a GenAI layer: implemented RAG using Sentence Transformer embeddings and ChromaDB to retrieve relevant business context, and built an LLM tool-calling workflow where the LLM triggers a Python backend to execute SQL against Redshift and return live, business-friendly answers via a natural language chatbot.
- Scheduled and monitored production jobs with Databricks Workflows; provided production support, troubleshooting pipeline failures and delivering business enhancement requests.
- Managed source control and deployments using Git, GitHub, and Azure DevOps within an Agile Scrum team.

TECHNICAL SKILLS

Programming Languages: Python, SQL, PySpark

Big Data & Data Engineering: Apache Spark, Databricks, Delta Lake, Medallion Architecture

Cloud Technologies: AWS S3, Amazon Redshift, Amazon Athena

Databases: MySQL

ETL & Data Engineering: ETL/ELT, Batch Processing, Data Pipelines, Data Transformation, Data Cleaning, Data Validation, Schema Validation, Data Quality

Scheduling & DevOps: Databricks Workflows, Azure DevOps, Git, GitHub

Analytics: Power BI, Window Functions, Advanced SQL Reporting

GenAI & AI Engineering: RAG, Vector Databases (ChromaDB), Sentence Transformers/Embeddings, Tool Calling, LLM Integration, Prompt Engineering, Natural Language Querying

Professional Practices: Production Support, Monitoring & Troubleshooting, Performance Optimization, Agile Scrum

EDUCATION

Bachelor of Computer Applications (BCA)

2025

SRTM University, Nanded, Maharashtra

CERTIFICATIONS

NPTEL - Python for Data Science, IIT Madras (Silver Medalist)

Jan-Feb 2025